



CSE 428

Milestone-2

Aabrar Islam - 20101361

Ayen Aziza Haque - 20301487

Simin Waliza - 20101401

Zahin Shabab - 20101165

Shahriar Azad Frahim- 20101223



Augmentation

Types of Augmentations used:

1. Rotation(left and right)
2. Flipping(left and right)
3. Zooming(in and out)
4. Changing contrast
5. Skewing

Augmentor

```
[ ] def augmentor(input_dir, output_dir, desired_count):  
    p = Augmentor.Pipeline(input_dir)  
    p.rotate(probability=0.3, max_left_rotation=10, max_right_rotation=10)  
    p.flip_left_right(probability=0.38)  
    p.zoom(probability=0.2, min_factor=1.1, max_factor=1.5)  
    p.skew(probability=0.3, magnitude=0.3)  
    p.random_contrast(probability=0.3, min_factor=0.8, max_factor=1.2)  
    original_count = len(p.augmentor_images)  
    images_needed = desired_count - original_count  
    while images_needed > 0:  
        p.sample(1)  
        images_needed -= 1  
    for root, _, files in os.walk(input_dir):  
        for filename in files:  
            source_path = os.path.join(root, filename)  
            if filename.startswith("output_"):  
                destination_path = os.path.join(output_dir, filename)  
                if not os.path.exists(destination_path):  
                    shutil.move(source_path, destination_path)  
            else:  
                destination_path = os.path.join(output_dir, filename)  
                if not os.path.exists(destination_path):  
                    shutil.copy(source_path, destination_path)  
    parent_folder = '/content/drive/MyDrive/428/Group_2/test'  
    output_parent_folder = '/content/drive/MyDrive/428/Group_2/aug_test'
```

Augmented Train Dataset

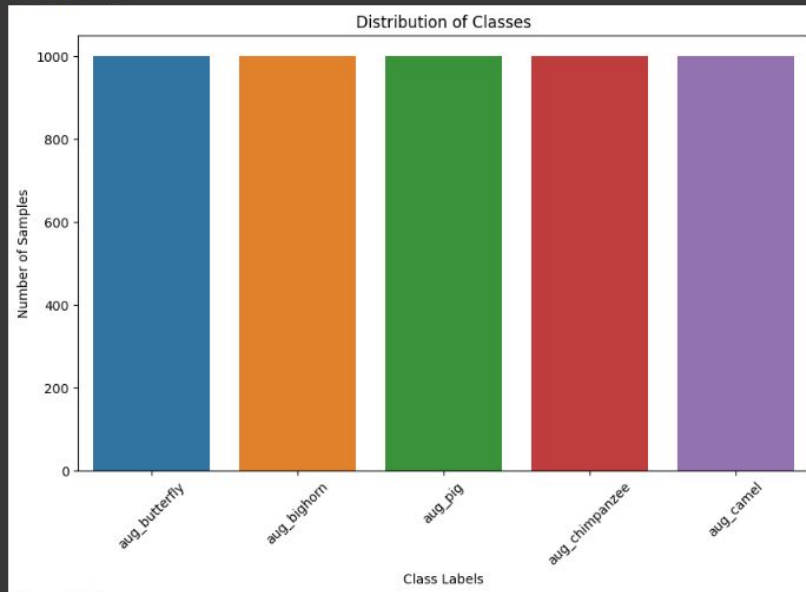
- 5 classes
- Train(5000)
 - Bighorn=1000
 - Butterfly=1000
 - Camel=1000
 - Chimpanzee=1000
 - Pig=1000

```
Class aug_butterfly: 1000 images
Class aug_bighorn: 1000 images
Class aug_pig: 1000 images
Class aug_chimpanzee: 1000 images
Class aug_camel: 1000 images
```

```
Total Images: 5000
```

```
Visualization:
```

```
=====
```



```
Balance Check:
```

```
=====
```

```
The distribution is balanced.
```

Augmented Test Dataset

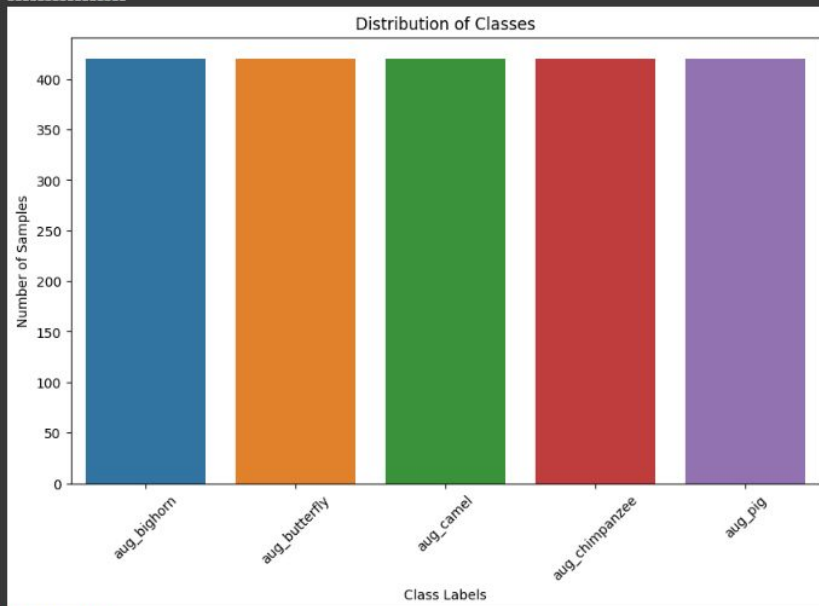
- 5 classes
- Test(2100)
 - Bighorn=420
 - Butterfly=420
 - Camel=420
 - Chimpanzee=420
 - Pig=420

```
Class aug_bighorn: 420 images
Class aug_butterfly: 420 images
Class aug_camel: 420 images
Class aug_chimpanzee: 420 images
Class aug_pig: 420 images
```

Total Images: 2100

Visualization:

=====



Balance Check:

=====

The distribution is balanced.

Performance Metric Formulae

1. **Sensitivity:** $\text{True Positive} / (\text{True Positive} + \text{False Negative})$
2. **Specificity :** $\text{True Negative} / (\text{True Negative} + \text{False Positive})$
3. **Positive Predictive Value(PPV):** $\text{True Positive} / (\text{True Positive} + \text{False Positive})$
4. **Negative Predictive Value(NPV):** $\text{True Negative} / (\text{True Negative} + \text{False Negative})$
5. **F Score:** $(2 * \text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$

```
def evaluate(model, X, y):
    y_pred_probs = model.predict(X)
    y_pred = np.argmax(y_pred_probs, axis=1)
    y_true = np.argmax(y, axis=1)
    accuracy = accuracy_score(y_true, y_pred)
    conf_matrix = confusion_matrix(y_true, y_pred)
    f1 = f1_score(y_true, y_pred, average='weighted')
    sensitivity, specificity, PPV, NPV = [], [], [], []
    for i in range(len(conf_matrix)):
        TP = conf_matrix[i, i]
        FN = conf_matrix[i, :].sum() - TP
        FP = conf_matrix[:, i].sum() - TP
        TN = conf_matrix.sum() - (TP + FN + FP)

        sensitivity.append(TP / (TP + FN) if (TP + FN) != 0 else 0)
        specificity.append(TN / (TN + FP) if (TN + FP) != 0 else 0)
        PPV.append(TP / (TP + FP) if (TP + FP) != 0 else 0)
        NPV.append(TN / (TN + FN) if (TN + FN) != 0 else 0)

    print("Accuracy:", accuracy)
    print("Sensitivity (per class):", sensitivity)
    print("Specificity (per class):", specificity)
    print("Positive Predictive Value (PPV) (per class):", PPV)
    print("Negative Predictive Value (NPV) (per class):", NPV)
    print("F1 Score:", f1)
    plot_confusion_matrix(conf_matrix)
```

Resnet50

```
def identity_block(X, f, filters):
    F1, F2, F3 = filters

    X_shortcut = X

    X = layers.Conv2D(F1, (1, 1), strides=(1, 1), padding='valid')(X)
    X = layers.BatchNormalization(axis=3)(X)
    X = layers.Activation('relu')(X)

    X = layers.Conv2D(F2, (f, f), strides=(1, 1), padding='same')(X)
    X = layers.BatchNormalization(axis=3)(X)
    X = layers.Activation('relu')(X)

    X = layers.Conv2D(F3, (1, 1), strides=(1, 1), padding='valid')(X)
    X = layers.BatchNormalization(axis=3)(X)

    X = layers.add([X, X_shortcut])
    X = layers.Activation('relu')(X)

    return X
```

```
def convolutional_block(X, f, filters, s=2):
    F1, F2, F3 = filters

    X_shortcut = X

    X = layers.Conv2D(F1, (1, 1), strides=(s, s), padding='valid')(X)
    X = layers.BatchNormalization(axis=3)(X)
    X = layers.Activation('relu')(X)

    X = layers.Conv2D(F2, (f, f), strides=(1, 1), padding='same')(X)
    X = layers.BatchNormalization(axis=3)(X)
    X = layers.Activation('relu')(X)

    X = layers.Conv2D(F3, (1, 1), strides=(1, 1), padding='valid')(X)
    X = layers.BatchNormalization(axis=3)(X)

    X_shortcut = layers.Conv2D(F3, (1, 1), strides=(s, s), padding='valid')(X_shortcut)
    X_shortcut = layers.BatchNormalization(axis=3)(X_shortcut)

    X = layers.add([X, X_shortcut])
    X = layers.Activation('relu')(X)

    return X
```

```

def ResNet50(input_shape=(64, 64, 3), classes=5):
    X_input = layers.Input(input_shape)

    X = layers.ZeroPadding2D((3, 3))(X_input)

    X = layers.Conv2D(64, (7, 7), strides=(2, 2))(X)
    X = layers.BatchNormalization(axis=3)(X)
    X = layers.Activation('relu')(X)
    X = layers.MaxPooling2D((3, 3), strides=(2, 2))(X)

    X = convolutional_block(X, f=3, filters=[64, 64, 256], s=1)
    X = identity_block(X, 3, [64, 64, 256])
    X = identity_block(X, 3, [64, 64, 256])

    X = convolutional_block(X, f=3, filters=[128, 128, 512], s=2)
    X = identity_block(X, 3, [128, 128, 512])
    X = identity_block(X, 3, [128, 128, 512])
    X = identity_block(X, 3, [128, 128, 512])

    X = convolutional_block(X, f=3, filters=[256, 256, 1024], s=2)
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])

    X = convolutional_block(X, f=3, filters=[512, 512, 2048], s=2)
    X = identity_block(X, 3, [512, 512, 2048])
    X = identity_block(X, 3, [512, 512, 2048])

```

```

X = layers.AveragePooling2D((2, 2), name='avg_pool')(X)

X = layers.Flatten()(X)
X = layers.Dense(classes, activation='softmax', name='fc' + str(classes))(X)

model = Model(inputs=X_input, outputs=X, name='ResNet50')

return model

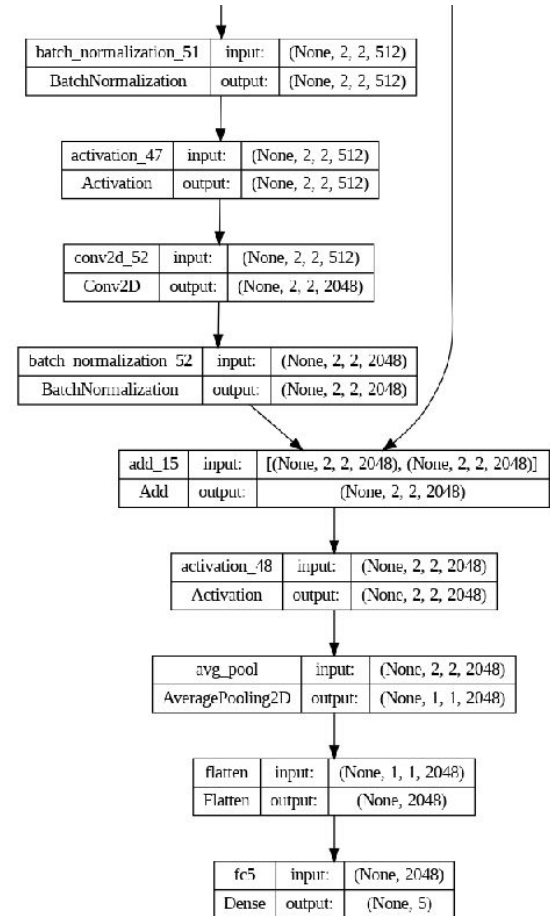
```

```

model = Sequential()
model = ResNet50(input_shape= (64,64,3), classes=num_classes)
model.summary()

```

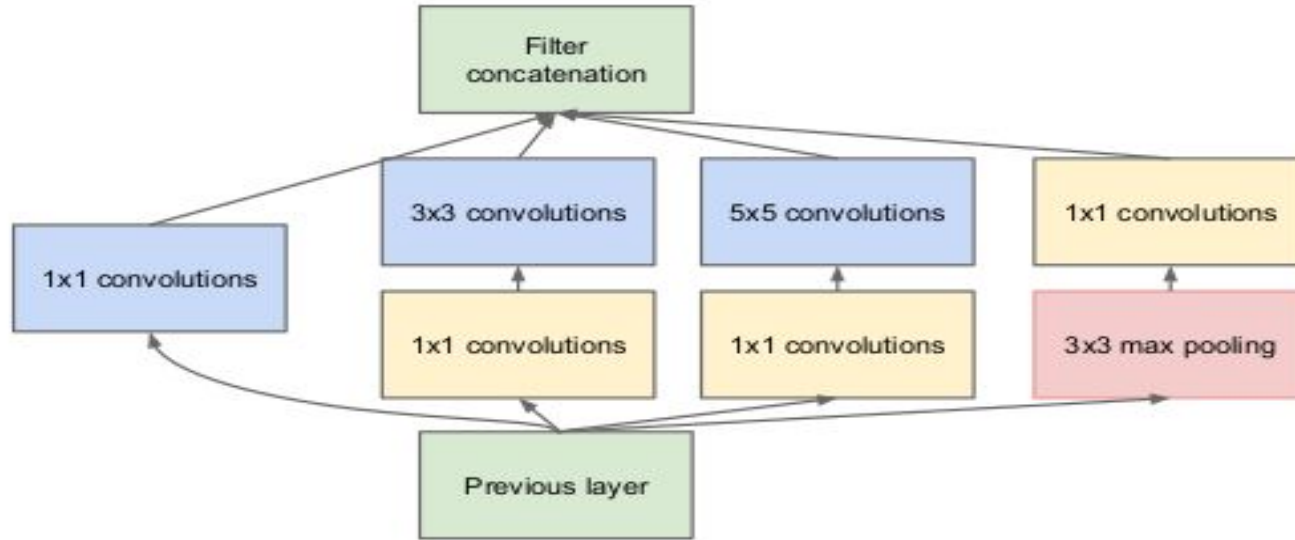
Structure of Resnet-50



Inception Net

```
1 def inception_module(x, f1, f2, f3):
2     # 1x1 conv
3     conv1 = keras.layers.Conv2D(f1, (1,1), padding='same', activation='relu')(x)
4     # 3x3 conv
5     conv3 = keras.layers.Conv2D(f2, (3,3), padding='same', activation='relu')(x)
6     # 5x5 conv
7     conv5 = keras.layers.Conv2D(f3, (5,5), padding='same', activation='relu')(x)
8     # 3x3 max pooling
9     pool = keras.layers.MaxPooling2D((3,3), strides=(1,1), padding='same')(x)
10    # concatenate filters
11    out = keras.layers.merge.concatenate([conv1, conv3, conv5, pool])
12    return out
```

Structure of Inception Net



Inception Layer

Custom CNN

```
model = models.Sequential()

model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Conv2D(256, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(512, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))

model.add(layers.Dense(5, activation='softmax'))

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(x_train, y_train_one_hot, epochs=10, validation_data=(x_test, y_test_one_hot))
```

Structure of Custom CNN

conv2d_input	input:	[(None, 64, 64, 3)]
InputLayer		
float32	output:	[(None, 64, 64, 3)]

conv2d	input:	(None, 64, 64, 3)
Conv2D		
float32	output:	(None, 62, 62, 32)

max_pooling2d	input:	(None, 62, 62, 32)
MaxPooling2D		
float32	output:	(None, 31, 31, 32)

conv2d_1	input:	(None, 31, 31, 32)
Conv2D		
float32	output:	(None, 29, 29, 64)



max_pooling2d_1	input:	(None, 29, 29, 64)
MaxPooling2D		
float32	output:	(None, 14, 14, 64)

conv2d_2	input:	(None, 14, 14, 64)
Conv2D		
float32	output:	(None, 12, 12, 128)

max_pooling2d_2	input:	(None, 12, 12, 128)
MaxPooling2D		
float32	output:	(None, 6, 6, 128)

conv2d_3	input:	(None, 6, 6, 128)
Conv2D		
float32	output:	(None, 6, 6, 256)



max_pooling2d_3	input:	(None, 6, 6, 256)
MaxPooling2D	output:	(None, 3, 3, 256)
float32		

conv2d_4	input:	(None, 3, 3, 256)
Conv2D	output:	(None, 3, 3, 512)
float32		

max_pooling2d_4	input:	(None, 3, 3, 512)
MaxPooling2D	output:	(None, 1, 1, 512)
float32		

flatten	input:	(None, 1, 1, 512)
Flatten	output:	(None, 512)
float32		

dense	input:	(None, 512)
Dense	output:	(None, 512)
float32		

dense_1	input:	(None, 512)
Dense	output:	(None, 5)
float32		

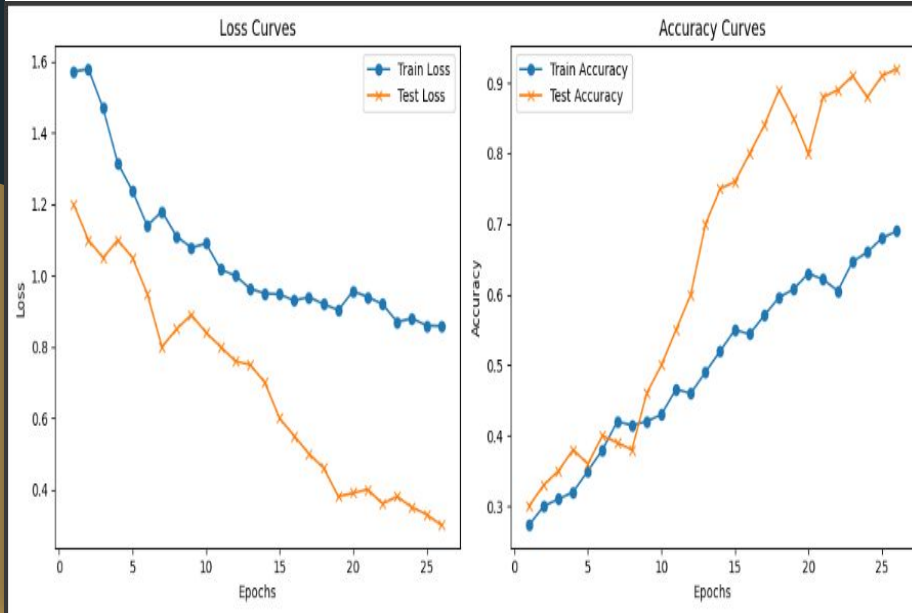
Better Model

The Resnet-50 model is the better model compared to the Inception Net model in this case.


- Resnet-50 accuracy 87.74% and Inception Net accuracy 67%.
- Easier to train due to less complex architecture.
- Consistency and more deeper layers.

Train and Test Curve

Inception Model



Resnet-50 Model



Misclassified Test Images

True: camel
Predicted: bighorn



True: bighorn
Predicted: chimpanzee



True: bighorn
Predicted: chimpanzee



True: pig
Predicted: chimpanzee



True: pig
Predicted: butterfly



True: bighorn
Predicted: camel



True: butterfly
Predicted: camel



True: pig
Predicted: bighorn



True: chimpanzee
Predicted: bighorn



True: pig
Predicted: camel




True: pig
Predicted: butterfly



True: chimpanzee
Predicted: bighorn



True: bighorn
Predicted: chimpanzee



True: pig
Predicted: chimpanzee



True: bighorn
Predicted: camel



```
y_pred = model.predict(x_test)
y_pred_classes = y_pred.argmax(axis=1)
misclassified_indices = []
for i, (true_label, predicted_label) in enumerate(zip(y_test, y_pred_classes)):
    if true_label != predicted_label:
        misclassified_indices.append(i)
random_indices = random.sample(misclassified_indices, min(20, len(misclassified_indices)))
fig, axes = plt.subplots(3, 5, figsize=(15, 8))
for i, idx in enumerate(random_indices):
    if i >= 15:
        break
    image = X_test[idx].reshape(64, 64, 3)
    true_label = y_test[idx]
    predicted_label = y_pred_classes[idx]
    true_class_name = class_names[true_label]
    predicted_class_name = class_names[predicted_label]
    ax = axes[i // 5, i % 5]
    ax.imshow(image)
    ax.set_title(f"True: {true_class_name}\nPredicted: {predicted_class_name}", color='red')
    ax.axis('off')
plt.tight_layout()
plt.show()
```




Updated Milestone-1

Aabrar Islam - 20101361

Ayen Aziza Haque - 20301487

Simin Waliza - 20101401

Zahin Shabab - 20101165

Shahriar Azad Frahim- 20101223



Dataset

- Size of the dataset is 2414
- Test set = 250
- Train set = 2164

- Separate batch for test set
- Separate batch for train set
- Validation part has to be done manually

Classes in the dataset

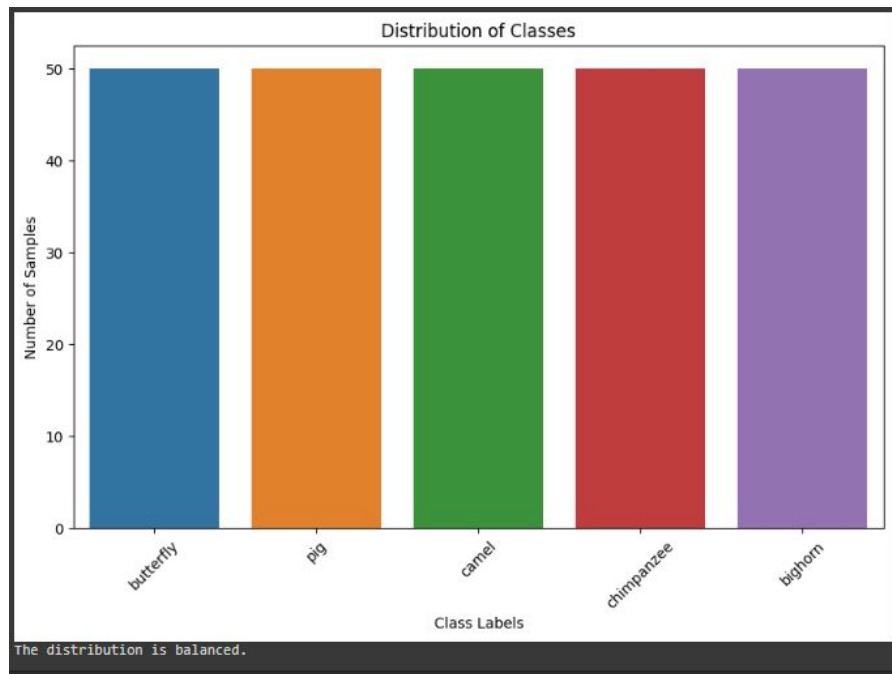
5 classes for both test and train sets are:

1. Bighorn
2. Butterfly
3. Camel
4. Chimpanzee
5. Pig.

```
Found 2164 images belonging to 5 classes.  
Found 250 images belonging to 5 classes.  
Class names: ['bighorn', 'butterfly', 'camel', 'chimpanzee', 'pig']
```

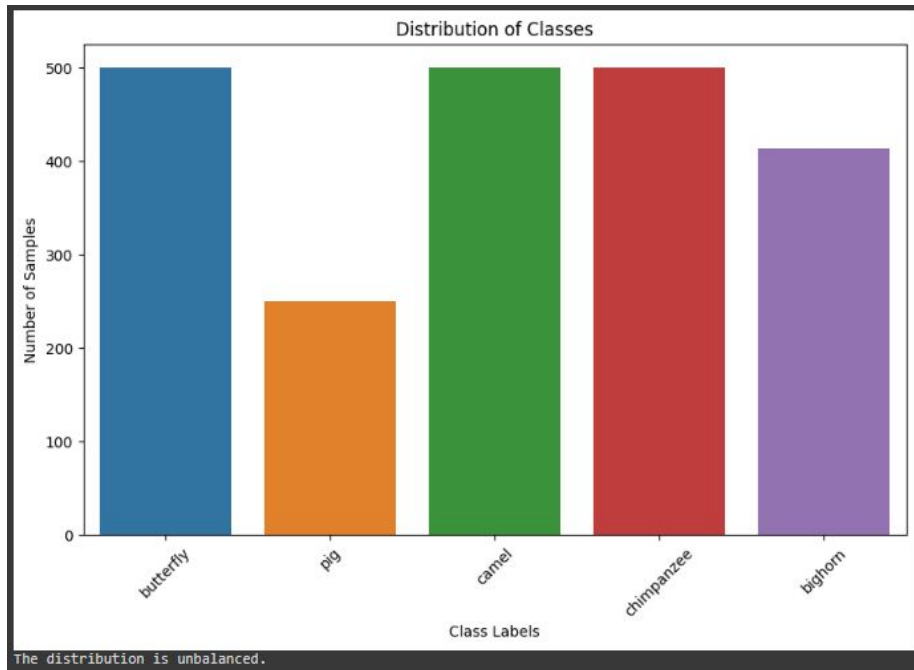
Distribution

- 5 classes
- Test(250)
 - Bighorn=50
 - Butterfly=50
 - Camel=50
 - Chimpanzee=50
 - Pig=50



Distribution

- 5 classes
- Train(2164)
 - Bighorn=414
 - Butterfly=500
 - Camel=500
 - Chimpanzee=500
 - Pig=250



Issues when dataset is unbalanced

1. **Biased Models:** Model may become biased toward the majority class, as it has more examples to learn from. This can result in poor performance on the minority class.
2. **Poor Generalization:** Unbalanced datasets can lead to poor generalization to new, unseen data, especially for the minority class. The model may simply predict the majority class most of the time, as it's more likely to be correct according to the training data distribution.
3. **Misleading Evaluation Metrics:** Common evaluation metrics like accuracy can be misleading when dealing with unbalanced datasets. A model that predicts the majority class all the time might have a high accuracy, but it's not useful if the minority class is of interest.

Tackling Unbalanced Dataset

1. **Resampling:**

- a. **Oversampling:** By producing duplicate or fake samples, increase the number of examples in the minority class.
- b. **Undersampling:** Reduce the number of instances in the majority class by deleting samples at random.

2. **Data Augmentation:** For image or text data, you can create new training examples by applying transformations or adding noise to existing samples. This can help balance the dataset.

3. **Synthetic Data Generation:** Use techniques like Synthetic Minority Over-sampling Technique (SMOTE) to generate synthetic examples for the minority class, based on the existing minority samples.

4. **Cost-sensitive Learning:** Assign different misclassification costs to different classes, emphasizing the importance of correctly classifying the minority class.
5. **Algorithm Selection:** Select techniques that are less sensitive to class imbalances, such as gradient boosting or support vector machines, which may be adjusted to place greater weight on minority class data.
6. **Change the Decision Threshold:** Adjust the categorization threshold to strike a balance between recall and precision based on your particular scenario. If the cost of false positives and false negatives differs, this can be especially helpful.
7. **Collect More Data:** In some cases, collecting more data for the minority class can be a solution if it's feasible.

Technique used to tackle unbalanced dataset

We chose the algorithm selection method to avoid or tackle the unbalanced dataset. Tree based algorithms are also performed in these scenarios and can also be boosted for unbalanced dataset


Algorithm Selection: Select techniques that are less sensitive to class imbalances, such as gradient boosting or support vector machines, which may be adjusted to place greater weight on minority class data.

Evaluation function

```
def evaluate(model, X, y):  
    y_pred = model.predict(X)  
    if len(y_pred.shape) > 1 and y_pred.shape[1] > 1:  
        y_pred = y_pred.argmax(axis=1)  
    accuracy = accuracy_score(y, y_pred)  
    conf_matrix = confusion_matrix(y, y_pred)  
    f1 = f1_score(y, y_pred, average='weighted')  
    sensitivity, specificity, PPV, NPV = [], [], [], []  
    for i in range(len(conf_matrix)):  
        TP = conf_matrix[i, i]  
        FN = conf_matrix[i, :].sum() - TP  
        FP = conf_matrix[:, i].sum() - TP  
        TN = conf_matrix.sum() - (TP + FN + FP)  
  
        sensitivity.append(TP / (TP + FN) if (TP + FN) != 0 else 0)  
        specificity.append(TN / (TN + FP) if (TN + FP) != 0 else 0)  
        PPV.append(TP / (TP + FP) if (TP + FP) != 0 else 0)  
        NPV.append(TN / (TN + FN) if (TN + FN) != 0 else 0)  
  
    print("Accuracy:", accuracy)  
    print("Sensitivity (per class):", sensitivity)  
    print("Specificity (per class):", specificity)  
    print("Positive Predictive Value (PPV) (per class):", PPV)  
    print("Negative Predictive Value (NPV) (per class):", NPV)  
    print("F1 Score:", f1)  
    plot_confusion_matrix(conf_matrix)
```


Logistic Regression(Train)

```
Accuracy: 1.0  
Sensitivity (per class): [1.0, 1.0, 1.0, 1.0, 1.0]  
Specificity (per class): [1.0, 1.0, 1.0, 1.0, 1.0]  
Positive Predictive Value (PPV) (per class): [1.0, 1.0, 1.0, 1.0, 1.0]  
Negative Predictive Value (NPV) (per class): [1.0, 1.0, 1.0, 1.0, 1.0]  
F1 Score: 1.0
```




Logistic Regression(Test)

Accuracy: 0.492
Sensitivity (per class): [0.26, 0.9, 0.5, 0.7, 0.1]
Specificity (per class): [0.895, 0.935, 0.785, 0.81, 0.94]
Positive Predictive Value (PPV) (per class): [0.38235294117647056, 0.7758620689655172, 0.36764705882352944, 0.4794520547945205, 0.29411764705882354]
Negative Predictive Value (NPV) (per class): [0.8287037037037037, 0.9739583333333334, 0.8626373626373627, 0.9152542372881356, 0.8068669527896996]
F1 Score: 0.45698907576333186



Neural Network

```
8/8 [=====] - 0s 4ms/step  
Accuracy: 0.576  
Sensitivity (per class): [0.3, 0.96, 0.68, 0.86, 0.08]  
Specificity (per class): [0.91, 0.95, 0.855, 0.78, 0.975]  
Positive Predictive Value (PPV) (per class): [0.45454545454545453, 0.8275862068965517, 0.5396825396825397, 0.4942528735632184, 0.4444444444444444]  
Negative Predictive Value (NPV) (per class): [0.8387096774193549, 0.9895833333333334, 0.9144385026737968, 0.9570552147239264, 0.8091286307053942]  
F1 score: 0.5230870060284397
```



Comparison between logistic regression and neural network

- Neural network is better than logistic regression in this case.